

# Understanding the Overheads of Launching CUDA Kernels

## Motivation

- ▶ Nvidia GPUs can run 10,000s of threads on independent SMs (Streaming Multi-processors)
  - Not ideal for device-wide barriers
- ▶ Method for device-wide barriers in GPUs
  - Software barriers (example in [1])
  - Implicit barriers: launching separate kernels (impacts performance)
- ▶ Alternative ways to achieve the same goal
  - Grid synchronization or multi-grid synchronization [2]
  - Higher performance might come from lower occupancy [3]
- ▶ Implicit barrier (additional kernels) vs. single kernel
- ▶ **Question:**
  - When not to launch an additional kernel?
  - What is the penalty of using different kinds of barriers in CUDA?

## Background

- ▶ Different kinds of kernel launch methods.
  - **Traditional Launch**
  - **Cooperative Launch** (CUDA 9)  
Introduced to support grid synchronization
  - **Cooperative Multi-Device Launch** (CUDA 9)  
Introduced to support multi-grid synchronization
- ▶ Sleep instruction: wait specific nanosecond in GPU kernel.

## Micro-benchmark

- ▶ Definition
  - **Kernel Latency:** Total latency to run kernels, start from CPU thread launching a thread, end at CPU thread noticing that the kernel is finished.
  - **Kernel Overhead:** Latency that is not related to kernel execution.
  - **Additional Latency:** Considering that CPU thread have just called a kernel launch function, additional latency is the additional latency to launch an additional kernel.
  - **CPU Launch Overhead:** Latency of CPU calling a launch function.
  - **Small Kernel:** Kernel execution time is not the main reason for additional latency.
  - **Larger Kernel:** Kernel execution time is the main reason for additional latency.

```
global__ void null_kernel_DEP()
{
    repeat10(asm volatile("nanosleep.u32 1000;"));
}

//example of launchfunction: traditional, cooperative, multi_device_cooperative
start(timer);
repeats(launchfunction(null_kernel_DEP, launchparameters));
cudaDeviceSynchronize();
stop(timer);
```

Latency that related to kernel execution here is 10 us (wait unit 1000 ns here)


time=5X(CPU Launch Overhead)

CPU thread wait for GPU kernel

Kernel Latency

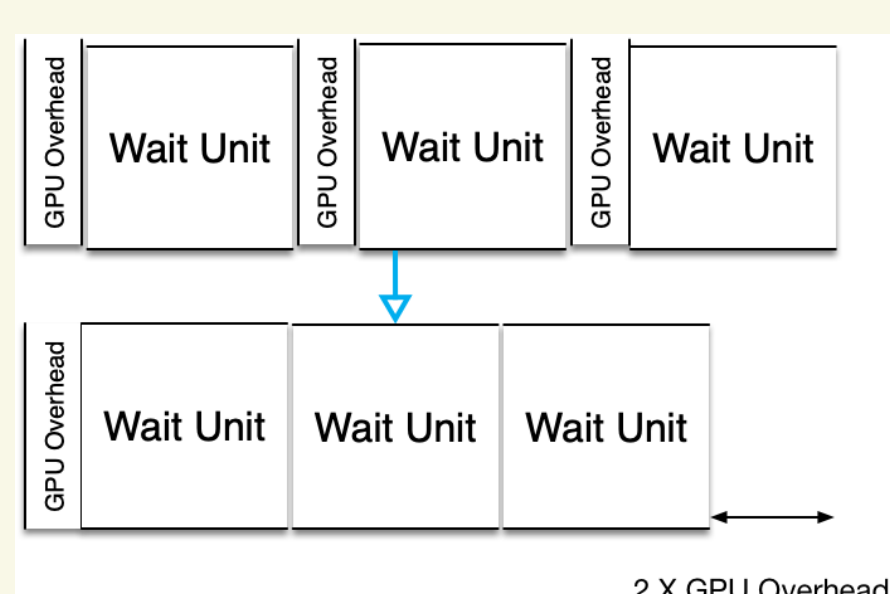
**Figure 1:** Sample code of micro-benchmark that call launch function 5 times, and repeats a wait unit (sleep 1000 ns) 10 times.

- ▶ Additional wait unit (sleep 1000 ns) do not increase any kernel overhead (Considering System Error)




**Figure 2:** Gradient of latency per wait unit (sleep 1000 ns) in a single kernel

- ▶ Test overhead in small kernels  
**Method:** Using null kernel (no code inside) to represent a Small Kernel
- ▶ Test overhead in large kernels  
**Method:** Using kernel fusion to unveil the overhead.



**Figure 3:** Using kernel fusion to test overhead hidden in kernel execution


## Launch Overhead in Small Kernels



**Figure 4:** Comparison of null kernel overhead using three different launch functions that employ different types of barriers (left), Cooperative Multi-Device Launch among different devices (right).

- ▶ **CPU Launch Overhead** is the main overhead in Small Kernel.

## Launch Overhead in Large Kernels




**Figure 5:** Comparison of Large Kernel Overhead among different launch functions (left), Cooperative Multi-Device Launch among different devices (right).

- ▶ CPU launch overhead is recorded to prove that it is not distinctive here. (the result is not as precise as the one in "Small Kernel" section)
- ▶ GPU execution overhead does exist.

## Other Overheads

- ▶ Empty kernel lasts about 8 us, still longer than the overheads we reported.



**Figure 6:** Comparison of different overheads in different launch functions

- ▶ **Other Overhead** is distinctive in single kernel. (Larger than the two kinds of overhead we reported)

## Conclusion

- ▶ Main overheads:
  - **Small Kernels:** CPU Launch Overhead
  - **Large Kernels:** GPU Execution Overhead
  - **Single Kernel:** Other Overhead
- ▶ Overhead of different launch functions
  - **Cooperative Multi-Device Launch > Cooperative Launch > Traditional Launch**
- ▶ Launch a new kernel when the performance improvement surpasses the overhead of a new kernel.

## References

- ▶ Shucai Xiao and Wu-chun Feng.  
Inter-block gpu communication via fast barrier synchronization.  
In *2010 IEEE International Symposium on Parallel & Distributed Processing (IPDPS)*, pages 1–12. IEEE, 2010.
- ▶ Cuda c programming guide, May 2019.
- ▶ Vasily Volkov.  
Better performance at lower occupancy.  
In *Proceedings of the GPU technology conference, GTC*, volume 10, page 16. San Jose, CA, 2010.

